

A PROPOSAL FOR THE REVENUE ORGANIZATION

Turn your SDR org into an AI-augmented growth engine

How MindPal's multi-agent workforce cuts research time by 90%+, drops lead response under 5 minutes, and reclaims ~\$700K/year in selling capacity — with a human approving every outbound message.

Prepared for: VP Sales / Chief Revenue Officer | Sales & Revenue Operations

July 2026

THE PROBLEM

Your reps aren't selling. They're doing data entry.

60%

of a rep's time goes to non-selling work —
writing
emails, data entry, admin and internal meetings

Salesforce State of Sales 2026 (n=4,050 reps)

28%

of an SDR's time specifically goes to researching
targets and sourcing contact info before ever
reaching out

Sopro, State of Prospecting

29 hrs

average time B2B companies take to respond to
a
lead — and that's only among the 36% who
ever reply at all

RevenueHero, 2024: 1,000 live demo requests tested

Bottom line: your best-comped, highest-potential people are spending 6 of every 10 working hours on tasks that require zero judgment — nearly a third of it just hunting for who to contact — and 64% of B2B companies never respond to an inbound lead at all.

THE COST

Here is what that's costing a 50-rep SDR organization

CAPACITY BURNED ON RESEARCH & ADMIN

\$1.4M / year

50 SDRs × \$100K fully-loaded cost (Bridge Group 2025 SDR Metrics Report) × 28% of time spent specifically on researching targets and sourcing contacts (Sopro).


That's the equivalent of 14 full-time reps doing lookup work instead of building pipeline.

THE REVENUE LEAK: SPEED

21x

more likely to qualify a lead if contacted within 5 min vs. 30 min

CONVERSION RATE BY RESPONSE SPEED



Sources: Lead Response Mgmt. study (Oldroyd et al., MIT/InsideSales); 2026 Speed-to-Lead Benchmark, ArtemisGTM (15M sales interactions analyzed); Sopro State of Prospecting

THE SOLUTION

MindPal: a coordinated team of AI agents, not another chatbot

MindPal is a self-serve, no-code platform: you (or your RevOps team) build the agents and connect the workflow yourselves — either by dragging nodes manually, or describing the process to Mindie, MindPal's AI workflow assistant, in plain language.

A

Specialist Agents

Each agent has one job — research, scoring, drafting — trained on your ICP, playbooks and brand voice via System Instructions and Knowledge Sources.

B

Workflows, Not One-Offs

Agents hand off work like an assembly line: Router, Loop and Orchestrator-Worker nodes replicate your actual sales process.

C

Human-in-the-Loop

A Human Input Node pauses the workflow for a rep to review, edit or approve before anything ships — nothing sends unsupervised.

D


Plugs Into Your Stack

Connects to your CRM and inbox via Composio (100+ apps incl. HubSpot) and MCP (Beta) — no rip-and-replace.

THE PRODUCT

This is the whole build surface

Not a mockup — MindPal's actual workflow canvas. Everything described on the previous slide is built here, by your team, without engineering.



Built by dragging — or by asking

Here Mindie was told “add a review step before delivery.” It inserted a Quality Check node between Write Draft and Deliver Output, using an Evaluator-Optimizer pattern.

Real branching, not a linear script

Route Task splits Research and Analyze Data to run in parallel, then rejoins them for the draft — the same shape as your actual sales process.

RevOps owns this surface

No code, no ticket, no deploy. The same canvas holds the Human Input Node that gates every message before it reaches a prospect.

The Research & Qualification Agent

BEFORE

SDR spends 15-20 min per account on manual research before the first touch

AFTER MINDPAL

Research Agent builds a full account brief in under a minute — SDR just reviews it

HOW THE WORKFLOW RUNS

- 1 Research Agent pulls firmographics, tech stack, funding/news signals and prior CRM history the moment a target account is queued.
- 2 Qualification Agent scores fit against your ICP and buying-signal rules, filtering out accounts that will never convert.
- 3 A Human Input Node hands the SDR a one-page brief and a recommended angle — the rep decides, the agent never sends anything alone.
- 4 Qualified accounts route automatically into the rep's queue in CRM, fully enriched — zero manual data entry.

The Speed-to-Lead Response Agent

BEFORE

Average inbound lead waits 29 hours
for a first human response — if one ever comes

AFTER MINDPAL

First personalized response goes out in
under 5 minutes, 24/7 — rep is looped in live

HOW THE WORKFLOW RUNS

- 1 The moment a lead fills a form or replies to an email, the Response Agent drafts a personalized reply using the account's context and intent signal.
- 2 A Router Node classifies urgency: high-intent leads get an instant escalation ping to the assigned rep's phone.
- 3 Human Input Node: for anything outside pre-approved templates, a rep approves or edits before it sends — speed without losing control.
- 4 After-hours and weekend leads get covered automatically instead of sitting in a queue until Monday.

The Quote-to-Proposal Agent

BEFORE

AE spends 2-4 hours assembling a custom quote, proposal and redline-ready contract

AFTER MINDPAL

Draft proposal, pricing and first-pass contract markup ready in minutes

HOW THE WORKFLOW RUNS

- 1 Orchestrator-Worker workflow pulls deal terms from CRM, pricing rules from your rate card, and past win/loss data for similar deals.
- 2 Sub-agents draft the proposal narrative, build the pricing table, and flag any non-standard contract clauses for legal review.
- 3 Human Input Node: Deal Desk or AE reviews and approves before anything reaches the customer — agents draft, humans decide.
- 4 Turnaround for a standard proposal drops from days to same-day, directly shortening your sales cycle.

THE VALUE, QUANTIFIED

Capacity redeployed. Revenue unlocked. Speed gained.

Every formula is shown in full — swap in your own headcount, pipeline and close rate. Bold rates are sourced (slides 2–3); reclaim/uplift assumptions are marked as estimates for you to pressure-test. Capacity and revenue are two separate mechanisms — deliberately not added together.

CAPACITY REDEPLOYED

~\$700K / yr

≈14 FTE-equivalent of selling time for a 50-rep org — not a headcount cut.

FORMULA

$50 \times \$100K \times 28\% \times 50\%$

*reps × fully-loaded cost/rep (Bridge Group) ×
research-time share (Sopro) × est. reclaim rate*

REVENUE UNLOCKED

+\$240–720K ARR / yr

On a \$24M example pipeline, 20% close rate.
Conservative → optimistic.

FORMULA

$\$24M \times 5\text{--}15\% \times 20\%$

*pipeline × est. conversion lift (5% conservative /
15% optimistic — see speed curve, slide 3) × close
rate*

SPEED GAINED

**20 min → <1 min
29 hrs → <5 min**

Per-account research time and lead response
time.

FORMULA

—

*pilot target, not a measured result — baselined
Day 1, verified at Day 90 (see rollout plan)*

The platform fee is not the decision driver here

MindPal is priced per seat and AI credits, self-serve — no procurement cycle required. Real published pricing, July 2026.

WHAT MINDPAL COSTS

Pilot (1 pod, 10-12 reps)

Advanced plan — \$179/mo

5 editor + 20 user seats, 30K AI credits/mo

Full rollout (50 reps)

Ultra plan — \$449/mo

Unlimited seats, 100K AI credits/mo

White-glove option

MindPal Managed — \$2,000/mo

Dedicated AI automation engineer builds it for you

WHAT IT'S MODELED TO CREATE (SEE FORMULAS, SLIDE 9)

+\$240–720K ARR

modeled revenue for a 50-rep org — plus ~\$700K/yr of selling capacity redeployed (separate mechanism, not added — see slide 9)

Ultra plan (\$449/mo = \$5,388/yr) is ~2% of even the conservative case. But the licence isn't the real cost — RevOps time and rep adoption are. Budget for those.

HOW IT WORKS

Agents do the work. A human always holds the send button.



Every workflow can include a Human Input Node at the step that matters most to you — first outbound message, quote value above a threshold, or contract terms. Agents draft and route; your reps stay the final decision-maker on anything that reaches a prospect.

Connects to your CRM and inbox via Composio (100+ apps incl. HubSpot) and MCP (Beta, via Zapier/Make/Composio servers) — no migration required.

Why not just use what you already have?

The three alternatives your team will raise in this meeting — and where each one stops.

A

vs. Your CRM's built-in AI

Agentforce and Breeze are strong inside their own walls — but they assume one vendor's data model and sales motion, and typically arrive as an enterprise add-on with a procurement cycle attached. MindPal sits on top of the stack you already run.

B

vs. Point tools

Enrichment tools, sequencers and AI SDRs each nail one step well. None of them own the handoff between steps — that glue stays manual. MindPal orchestrates the whole chain, with a human gate wherever you choose to put one.

C

vs. Building it in-house

Your RevOps team configures this in days, no code. The same orchestration on raw frameworks means engineering headcount up front — and permanent maintenance you own forever.

D

So the real question

Isn't "can AI do this?" — every option says yes. It's who runs all three use cases end-to-end, on the stack you already own, with a human on the send button — provable in 90 days for \$179/mo.

From pilot to full deployment in one quarter

Days 1-30

Pilot Setup

- Start Advanced plan (\$179/mo) for 1 pod (10-12 reps)
- Connect CRM and inbox via Composio
- Configure Research + Qualification agents on your ICP
- Baseline current response time & research time

Days 31-60

Live Pilot

- Launch Speed-to-Lead agent with a Human Input Node for review
- Weekly tuning with RevOps on scoring & templates
- Track response time, SQL volume, hours reclaimed

Days 61-90

Prove & Decide

- Deal Desk / Proposal agent added for pod's AEs
- Full ROI readout vs. baseline
- Decide to upgrade to Ultra plan (\$449/mo) and scale to all 50 reps

THE ASK

One pod. 90 days. A decision you can make today.

- 1 Plan**

Start the pilot on MindPal's Advanced plan — \$179/month, self-serve signup, no procurement cycle. Upgrade to Ultra (\$449/mo) only once the pilot proves out.
- 2 People**

One pod of 10-12 reps, a pod manager as sponsor, and ~3 hrs/week from RevOps to configure agents and connect tools via Composio.
- 3 Access**

Login credentials to connect CRM and inbox via Composio (OAuth) for the pilot pod — no IT migration required.
- 4 Success bar**

Agreed upfront: <5 min lead response, +20% SQL volume, 10+ hrs/week/rep reclaimed — measured at Day 90.

At \$179/month, this pilot costs less than a single team lunch — if we don't hit the numbers by Day 90, we simply don't renew.

The elevator pitch

Your SDRs spend 60% of their time on non-selling work — nearly a third of it just researching who to contact — while the average inbound lead waits 29 hours for a response, if one ever comes at all. On MindPal, your team builds a coordinated set of AI agents — with a Human Input Node holding every outbound message for review — that cuts research time by over 90%, drops response time under 5 minutes, and is modeled to redeploy roughly \$700K a year of selling capacity in a 50-rep org. At \$179–449 a month self-serve with no procurement cycle, this isn't a budget decision — it's a 90-day pilot with one pod to prove the numbers, so let's put one pod on it.

Let's put one pod on it.